Lecture-6a

Predicting Missing Links

Feb 20, 2024.

Reminder: Project proposal **hard**.
PSA

actual ⇒ $G = (V, E)$

observed ⇒ $G^* = (V, E^*)$, where $E^* \subseteq E$ $\left(E^* = f(E) \right)$

unconnected pairs $Y = V \times V - E^*$ ~ $\Theta(n^2)$ haystack

missing links $X = E - E^*$ ~ $O(n)$ needles.

missingness function.

baseline accuracy by guessing $O(\frac{1}{n})$

[so hard!]

Reasons we care:

- most # are incomplete! (ex: - observing ecosystem, - gene-gene interaction, - NSA / CIA)
- use correlations to marshall resources.
- compare methods or models. (cross-validation)

Facebook as eg.
1. recency
2. friends of friends.

Idea of a predictor is to optimize a score function.

$$\text{score}(i, j) : Y \to \mathbb{R}$$

Key message:

We can define a "spurious link prd", but it will [redacted]
(by changing the score function)
be easier than link prediction. (why?)

Ⓗ (n^2) to explore v.s. Ⓗ (n) to explore.

A baseline algm.

if $i, j \in Y$, $score(i, j) = Uniform(0, 1)$

↓

Independent of G^* !

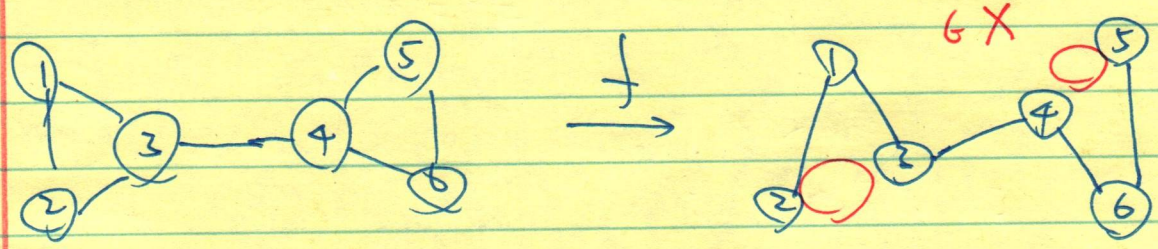

score table

i	j	$score(i, j)$
1	5	$r = Uniform(0, 1)$
1	4	r
1	6	r
$\in Y$	:	r
2	3	r
:	:	r
4	5	:

Topological predictors and local smoothness.

the idea

predict a missing link to occur where it would cluster with other (observed) edges

↓ define this as a predictor

Two ways to operationalize this idea:

1) Jaccard coefficient.

$$\text{Jaccard}(i,j) = \frac{|nb(i) \cap nb(j)|}{|nb(i) \cup nb(j)|}$$

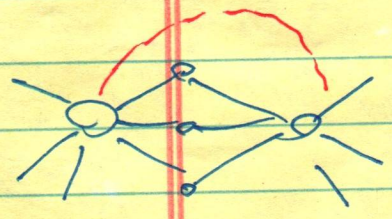

2) degree product $= k_i k_j$

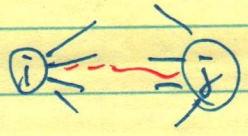

$$score(i,j) = k_i k_j + Uniform(0, \epsilon) \quad \epsilon \ll 1$$

$$scores(i,j) = Jaccard(i,j) + Uniform(0, \epsilon)$$

↑ break ties randomly

$\epsilon \ll \frac{1}{n-2}$

④

an example Jaccard :

DP :

another example :

(5)

ask ??

§ Measuring performance : the AUC

predicting missing links is a binary classifier. (reg. only TP / FP).

we can use Area Under the Curve (AUC)

↘ Receiver operating characteristics (ROC) curve.

nice properties

1) scale invariant (only cares about the ordering)
2) threshold invariant (general measure of distinguishability) relative

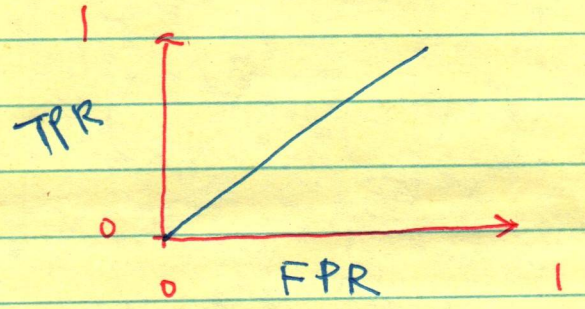

Mathematically

$$AUC = \Pr\left(score(TP) > score(TN)\right)$$

AUC = 1 ⇒ perfect distinguishability

= 1/2 ⇒ if I cannot tell a TP from TN.

In link prediction

$|y-x|$

TN: $\Theta(n^2)$

$|X|$

TP: $O(n)$

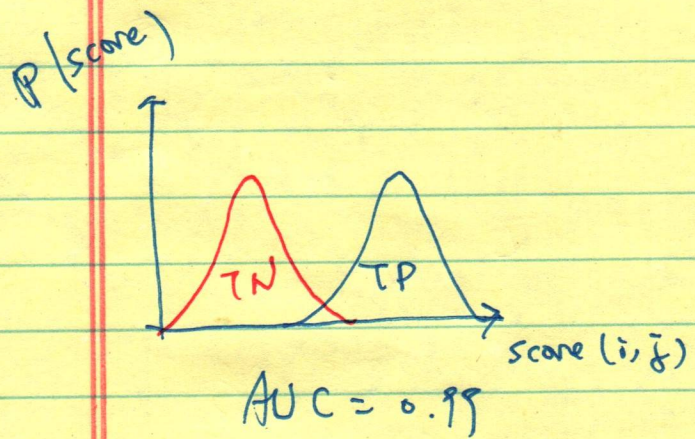

AUC = 0.99

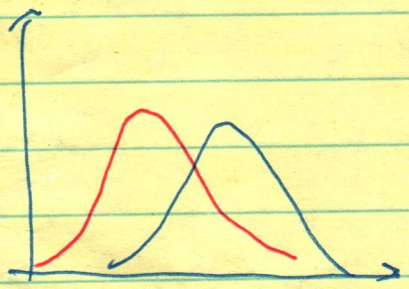

AUC = 0.76

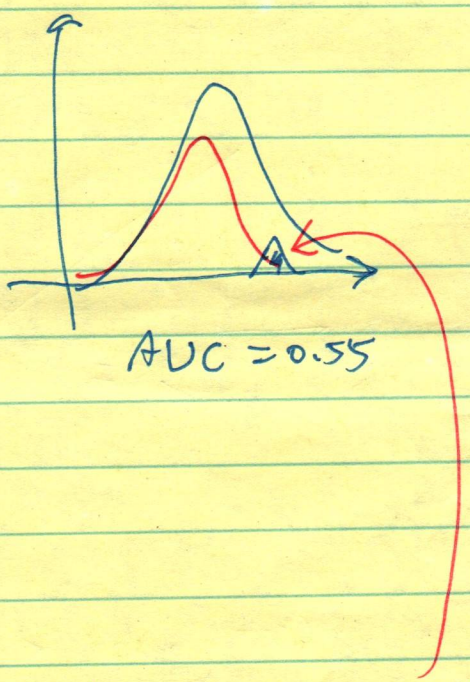

AUC = 0.55

AUC ⇒ useful to compare predictors (agnostically)

actually ...

if we want to see how much better, we need to look at "the curve"

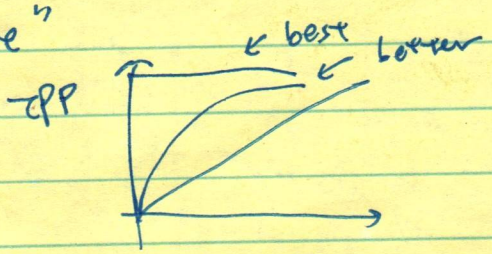

Wk-6b. Feb 22, 2024.

$$TPR(\ell) = \frac{1}{T} \sum_{k=1}^{\ell} \tau_k \quad ; \quad \tau_k = \begin{cases} 1 & \text{if it's missing a edge} \\ 0 & \text{otherwise} \end{cases}$$

(# TP)

$$FPR(\ell) = \frac{1}{F} \sum_{k=1}^{\ell} (1-\tau_k) \; ; \qquad FPR(0) \equiv 0.$$

(# FP)

$$AUC = \sum_{\ell=1}^{|T|} TPR(\ell) \times \left(FPR(\ell) - FPR(\ell-1) \right)$$

$TPR(\ell)$: y ; $FPR(\ell) - FPR(\ell-1)$: Δx

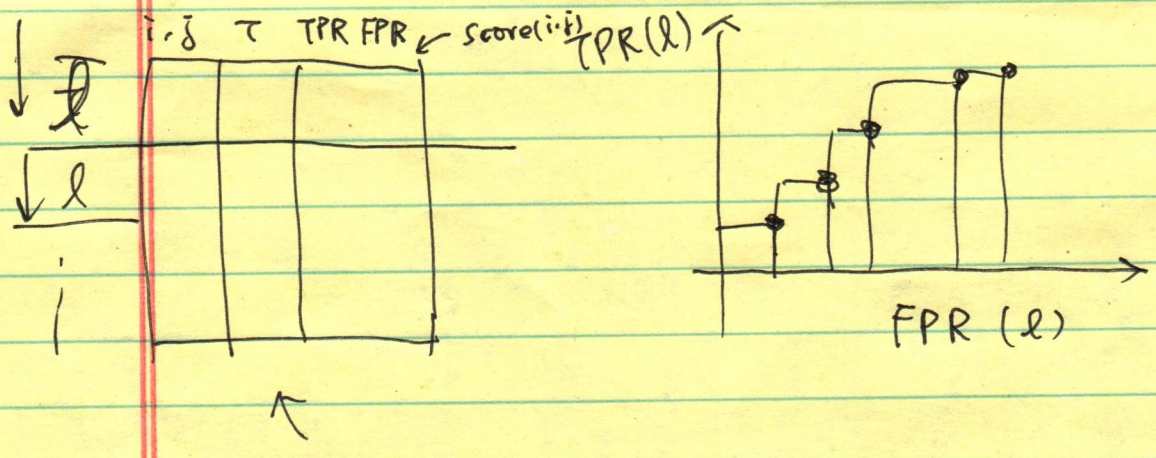

ROC coordinates

Example (See Clausets notes)

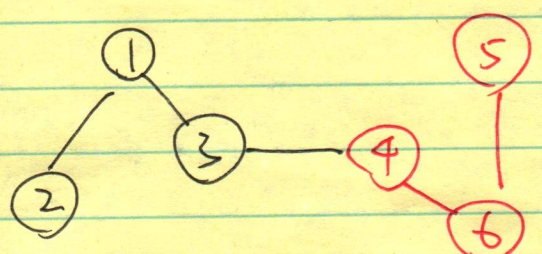

Observed $G^* = (V, E^*)$

Wk 6 - 6. Feb 22, 2024.

§ The world of missing link predictors.

Three main classes.

1) topological : score (i,j) use only G^* | local measure of structure around i,j

Jaccard, deg prod, $f(U_i, U_j)$ ← node-level statistic

(implicitly assumes an edge generating mechanism)

2) model-based

$\Pr(G|\theta) = \prod_{i<j} \Pr(i \to j \,|\, \theta)$ edges are conditionally indep

elaborate random graph models →

(assume some explicit edge generating mechanism)

Stochastic block model (SBM)

3) embeddings

$d(i,j)^{-1} \propto$ score (i,j)

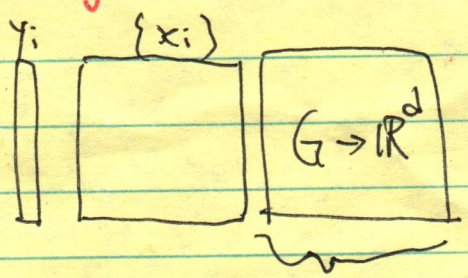

assumes a spatial homophily mechanism for edge generating.

But even more exciting are meta-learning solutions

(aka. ensemble learning)

→ key assumption: edges are the observations and the predictors

brief review

model selection via cross validation

observed data X

$\{x_i\}$

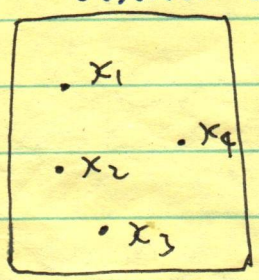

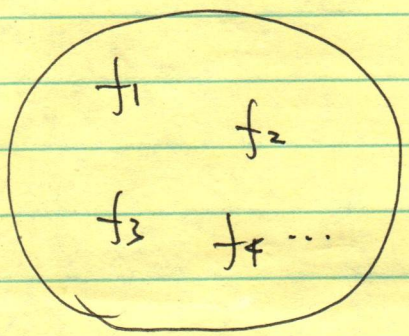

candidate models. F.

$\{f_n\}$

random forest, neural network, regression, etc.

* Which f_k provides the best out-of-sample prediction accuracy on not-yet-observed x_i ?

Solution:

① partition X in test and train data, randomly $X_{train} \cup X_{test}$

② learn apply each f_k to $X_{train} \mapsto \theta_k$

③ measure accuracy of each θ_k on $X_{test} \rightarrow \varphi_k$

④ Select $\underset{k}{\text{argmax}}\ \varphi_k \rightarrow$ best model

" winner take all approach "

"learning w ensembles"

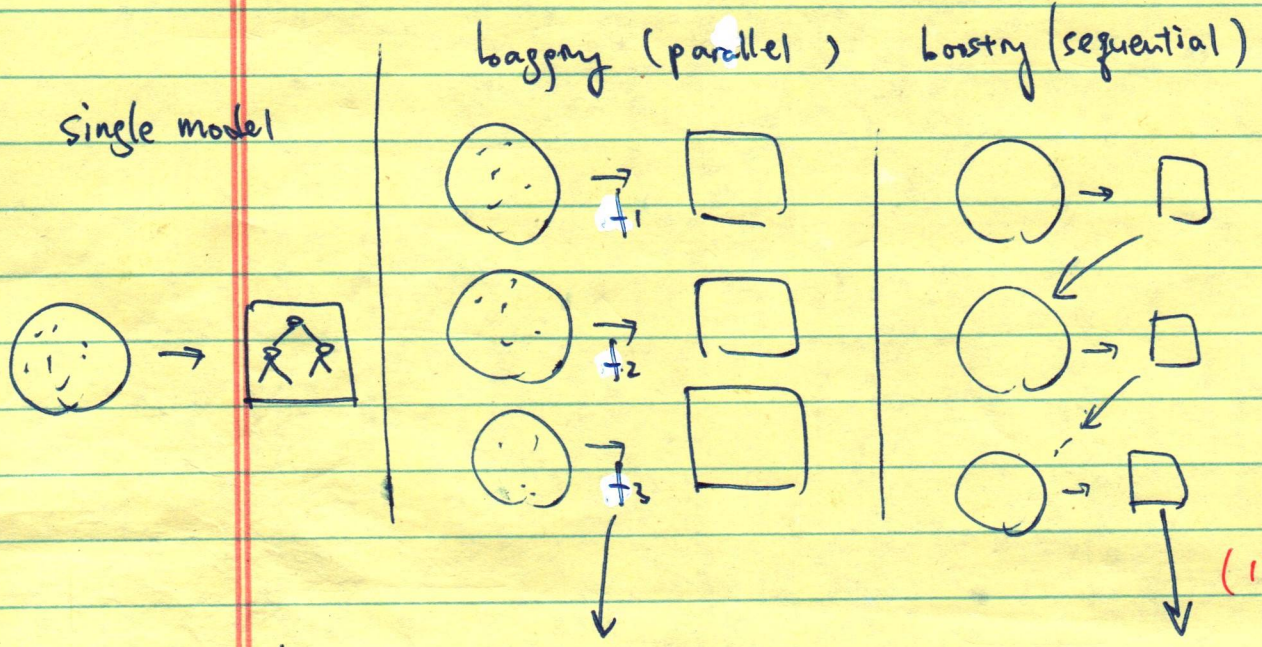

bootstrap aggregating (1994)
(special case of model averaging)

simulate noise in data

seq. produce "weak predictors"
[best single feature model]
up weight hardest remaining X_i
linearly combine weak predictors

AdaBoost
(main algm people use)

(5)

Stacking (1992)

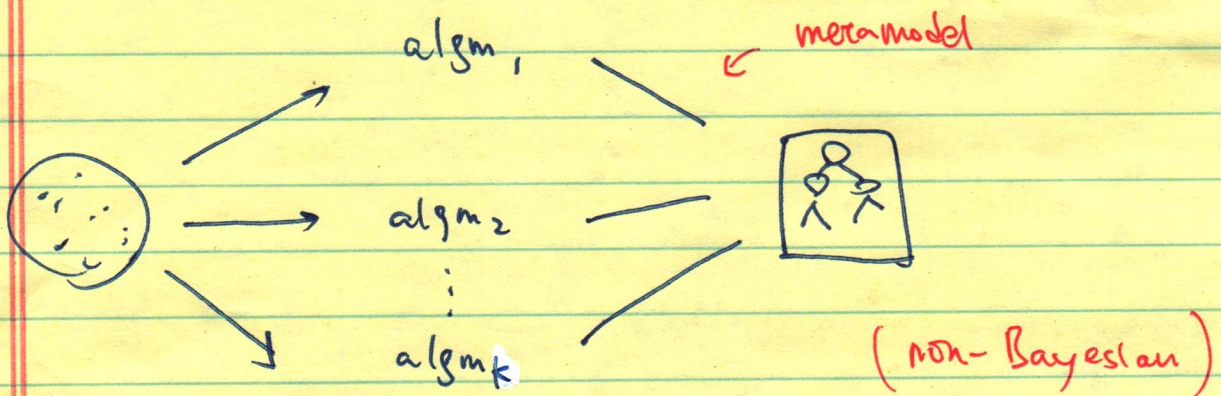

train many strong or weak predictors

learn a model of which predictor is best for a given type of input.

Delphi method
→ constructing a predictive distribution

not model selection or model averaging, but model combination.

(6)

35 min — start the slides.

predicted label \ true label	①	②
①	TN	FN
②	FP	TP

② → positive.

$$\text{recall} = \frac{TP}{TP + FN}$$

$$\text{precision} = \frac{TP}{TP + FP}$$

⇒ 203 features v. 25-30 features.

(7)

On real data :

1) inherently harder.
2) overfit — driven by algm.

~80% social
remaining (bio / info)
(ease (others)

model class (v.s.) all data (No-free lunch)